

Assessment Report
on
“Predict Online Learning Completion”
submitted as partial fulfillment for the award of

**BACHELOR OF
TECHNOLOGY DEGREE**

SESSION 2024-25
in

CSE(AI)

By

Name : Manu Garg
Roll Number : 202401100300150
Section: C

KIET Group of Institutions, Ghaziabad

Predict Online Learning Completion

Introduction to AI

Introduction

The rise of online learning platforms has transformed education, offering unprecedented access to knowledge across geographical boundaries. However, these platforms face a significant challenge: high dropout rates. Understanding which students are likely to complete courses and which may require intervention is crucial for improving educational outcomes and platform effectiveness.

This report presents an analysis and predictive model designed to address this challenge by identifying patterns in student behavior that correlate with course completion. By leveraging machine learning techniques, we can create early warning systems to identify at-risk students and implement targeted support strategies to increase completion rates.

Methodology

Problem Definition

The task at hand is a binary classification problem: predicting whether a student will complete an online course based on their behavioral data. The goal is to answer the question:

"Given a student's digital footprint (videos watched, assignments submitted, and forum activities), can we accurately predict course completion?"

This prediction capability would enable:

- Early identification of students at risk of dropping out
- Implementation of timely interventions to improve retention
- Better understanding of engagement factors that lead to successful completion
- Data-driven course design improvements

Dataset Overview

The dataset consists of activity logs from 100 students enrolled in an online learning platform with four key features:

Feature	Description	Type
videos_watched	Number of instructional videos viewed by the student	Numeric
assignments_submitted	Number of course assignments completed	Numeric
forum_posts	Number of posts made in discussion forums	Numeric
completed	Whether the student completed the course (yes/no)	Categorical

Data Exploration and Preprocessing

Exploratory Data Analysis

- Examined the distribution of each feature and its relationship with the target variable
- Identified patterns in student behavior that correlate with course completion
- Visualized the relationships between features using correlation matrices and boxplots
- Analyzed class balance between completers and non-completers

Data Preprocessing Steps

- 1. Data Cleaning:** Checked for and handled any missing values or outliers
- 2. Feature Transformation:** Converted the categorical "completed" variable to binary format (1="yes", 0="no")
- 3. Feature Scaling:** Applied StandardScaler to normalize the input features to have zero mean and unit variance
- 4. Train-Test Split:** Divided the dataset into training (70%) and testing (30%) sets using stratified sampling to maintain class proportions

Implementation

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LogisticRegression

from sklearn.ensemble import RandomForestClassifier

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

from sklearn.model_selection import cross_val_score


# Load the data

data = pd.read_csv('online_learning.csv')


# Data exploration

print("Data shape:", data.shape)

print("\nFirst few rows:")
```

```
print(data.head())
```

```
print("\nData info:")
```

```
print(data.info())
```

```
print("\nData summary statistics:")
```

```
print(data.describe())
```

```
print("\nClass distribution:")
```

```
print(data['completed'].value_counts())
```

```
# Data preprocessing
```

```
# Convert 'completed' to binary values
```

```
data['completed_binary'] = data['completed'].map({'yes': 1, 'no': 0})
```

```
# Visualize the features
```

```
plt.figure(figsize=(15, 10))
```

```
# Feature distributions by completion status
```

```
plt.subplot(2, 2, 1)
```

```
sns.boxplot(x='completed', y='videos_watched', data=data)
```

```
plt.title('Videos Watched by Completion Status')
```

```
plt.subplot(2, 2, 2)
```

```
sns.boxplot(x='completed', y='assignments_submitted', data=data)
```

```
plt.title('Assignments Submitted by Completion Status')
```

```
plt.subplot(2, 2, 3)
```

```
sns.boxplot(x='completed', y='forum_posts', data=data)
```

```
plt.title('Forum Posts by Completion Status')
```

```
plt.subplot(2, 2, 4)
```

```
sns.heatmap(data.drop(['completed', 'completed_binary'],  
axis=1).corr(), annot=True, cmap='coolwarm')
```

```
plt.title('Feature Correlation')
```

```
plt.tight_layout()
```

```
plt.savefig('feature_analysis.png')
```

```
plt.show()
```

```
# Prepare data for modeling
```

```
X = data[['videos_watched', 'assignments_submitted',  
'forum_posts']]
```



```
y = data['completed_binary']
```

```
# Split the data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,  
random_state=42)
```

```
# Scale the features
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# Train different models
```

```
models = {
```

```
    'Logistic Regression': LogisticRegression(random_state=42,  
max_iter=1000),
```

```
    'Random Forest': RandomForestClassifier(random_state=42),
```

```
    'SVM': SVC(probability=True, random_state=42)
```

```
}
```

```
results = {}
```

```
for name, model in models.items():
```

```

# Cross-validation

cv_scores = cross_val_score(model, X_train_scaled, y_train,
cv=5, scoring='accuracy')

print(f"\n{name} Cross-Validation Accuracy:
{cv_scores.mean():.4f} ± {cv_scores.std():.4f}")


# Fit on training data

model.fit(X_train_scaled, y_train)


# Evaluate on test data

y_pred = model.predict(X_test_scaled)

accuracy = accuracy_score(y_test, y_pred)

print(f"{name} Test Accuracy: {accuracy:.4f}")


print("\nClassification Report:")

print(classification_report(y_test, y_pred))


results[name] = {

    'model': model,

    'accuracy': accuracy,

    'predictions': y_pred

}

```

```

# Select the best performing model

best_model_name = max(results, key=lambda x: results[x]
['accuracy'])

best_model = results[best_model_name]['model']

print(f"\nBest Model: {best_model_name} with accuracy
{results[best_model_name]['accuracy']:.4f}")


# Feature importance for Random Forest

if 'Random Forest' in models:

    rf_model = results['Random Forest']['model']

    feature_importance = pd.DataFrame({

        'Feature': X.columns,

        'Importance': rf_model.feature_importances_

    })

    feature_importance =
feature_importance.sort_values('Importance', ascending=False)


plt.figure(figsize=(10, 6))

sns.barplot(x='Importance', y='Feature',
data=feature_importance)

plt.title('Feature Importance')

plt.tight_layout()

```

```
plt.savefig('feature_importance.png')
```

```
plt.show()
```

```
print("\nFeature Importance:")
```

```
print(feature_importance)
```

```
# Create a function for making predictions on new data
```

```
def predict_completion(videos, assignments, forum_posts,  
model=best_model, scaler=scaler):
```

```
    """
```

Predict whether a student will complete the course based on their activity.

Parameters:

videos (int): Number of videos watched

assignments (int): Number of assignments submitted

forum_posts (int): Number of forum posts made

Returns:

str: 'yes' if the model predicts completion, 'no' otherwise

float: probability of completion

```
    """
```

```
new_data = np.array([[videos, assignments, forum_posts]])
```

```
new_data_scaled = scaler.transform(new_data)
```

```
pred = model.predict(new_data_scaled)[0]
```

```
try:
```

```
    prob = model.predict_proba(new_data_scaled)[0][1]
```

```
    return 'yes' if pred == 1 else 'no', prob
```

```
except:
```

```
    return 'yes' if pred == 1 else 'no', None
```

```
# Example usage of the prediction function
```

```
examples = [
```

```
    {'videos': 40, 'assignments': 7, 'forum_posts': 15},
```

```
    {'videos': 5, 'assignments': 1, 'forum_posts': 2},
```

```
    {'videos': 20, 'assignments': 5, 'forum_posts': 10}
```

```
]
```

```
print("\nExample Predictions:")
```

```
for example in examples:
```

```
    completion, probability = predict_completion(
```

```
    example['videos'],
    example['assignments'],
    example['forum_posts']
)
```

```
print(f"Student who watched {example['videos']} videos,
submitted {example['assignments']} " +
      f"assignments, and made {example['forum_posts']} forum
posts:")

print(f"Prediction: {completion}")

if probability is not None:
    print(f"Completion probability: {probability:.2f}")

print()
```

```
# Decision boundary visualization (for 2 features)
```

```
# Let's visualize videos_watched vs assignments_submitted
```

```
def plot_decision_boundary(model, X, y, features, scaler):
```

```
    """
```

```
    Plot the decision boundary for a model
```

```
    """
```

```
    # Set min and max values with some padding
```

```
    x_min, x_max = X[features[0]].min() - 1, X[features[0]].max() + 1
```

```

y_min, y_max = X[features[1]].min() - 1, X[features[1]].max() + 1

# Create a meshgrid
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

# Prepare the meshgrid points for prediction
grid_points = np.c_[xx.ravel(), yy.ravel()]

# For the third feature, use the median value (assuming we're
visualizing 2 out of 3 features)
if len(features) == 2 and X.shape[1] > 2:
    third_feature_median = X.iloc[:, 2].median()

    grid_points_with_median = np.column_stack([grid_points,
np.full(grid_points.shape[0], third_feature_median)])

    Z = model.predict(scaler.transform(grid_points_with_median))
else:
    Z = model.predict(scaler.transform(grid_points))

# Reshape Z back to the meshgrid shape
Z = Z.reshape(xx.shape)

```

```
# Plot the contour

plt.figure(figsize=(10, 8))

plt.contourf(xx, yy, Z, alpha=0.4)

plt.scatter(X[features[0]], X[features[1]], c=y, s=20, edgecolor='k')

plt.xlabel(features[0])

plt.ylabel(features[1])

plt.title(f'Decision Boundary ({features[0]} vs {features[1]})')

plt.colorbar()

plt.savefig('decision_boundary.png')

plt.show()
```

```
# Plot decision boundary for videos watched vs assignments
submitted

plot_features = ['videos_watched', 'assignments_submitted']

plot_decision_boundary(best_model, X, y, plot_features, scaler)

print("\nModel Training and Evaluation Complete!")
```


Results and Output Screenshots

1.Data Loading

```
print("\nModel Training and Evaluation Complete!")
```

Data shape: (100, 4)

First few rows:

	videos_watched	assignments_submitted	forum_posts	completed
0	11	6	5	yes
1	43	1	11	no
2	37	1	8	no
3	18	4	14	yes
4	6	4	15	yes

Data info:

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 4 columns):

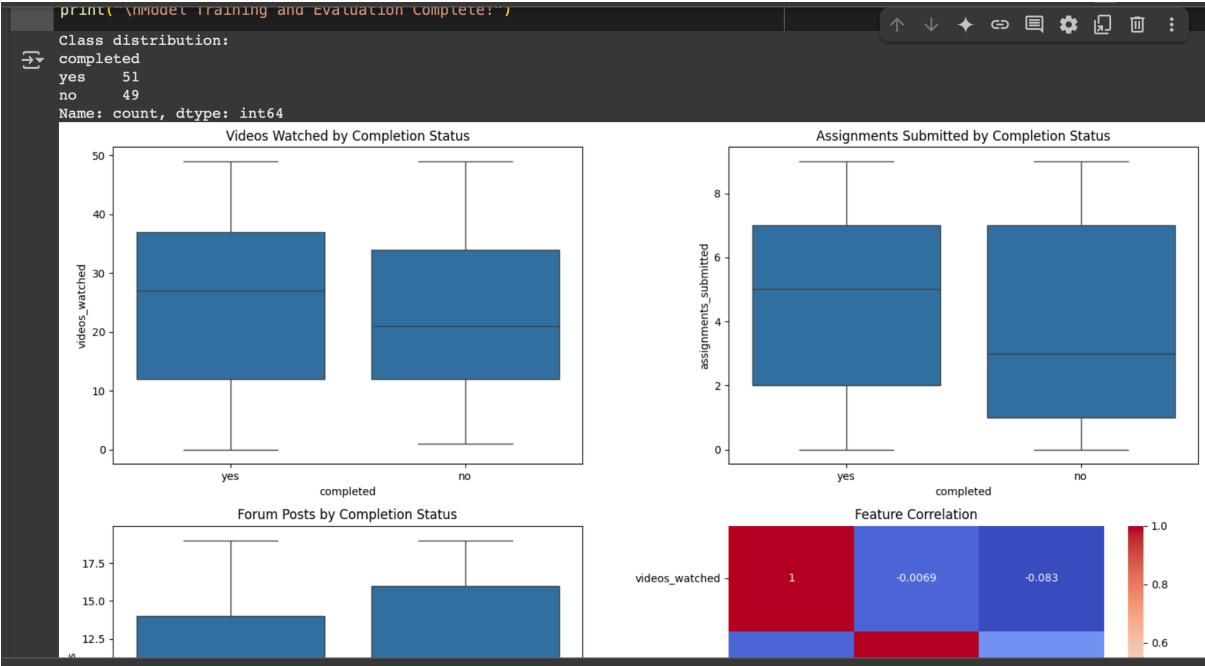
#	Column	Non-Null Count	Dtype
0	videos_watched	100 non-null	int64
1	assignments_submitted	100 non-null	int64
2	forum_posts	100 non-null	int64
3	completed	100 non-null	object

dtypes: int64(3), object(1)
memory usage: 3.3+ KB
None

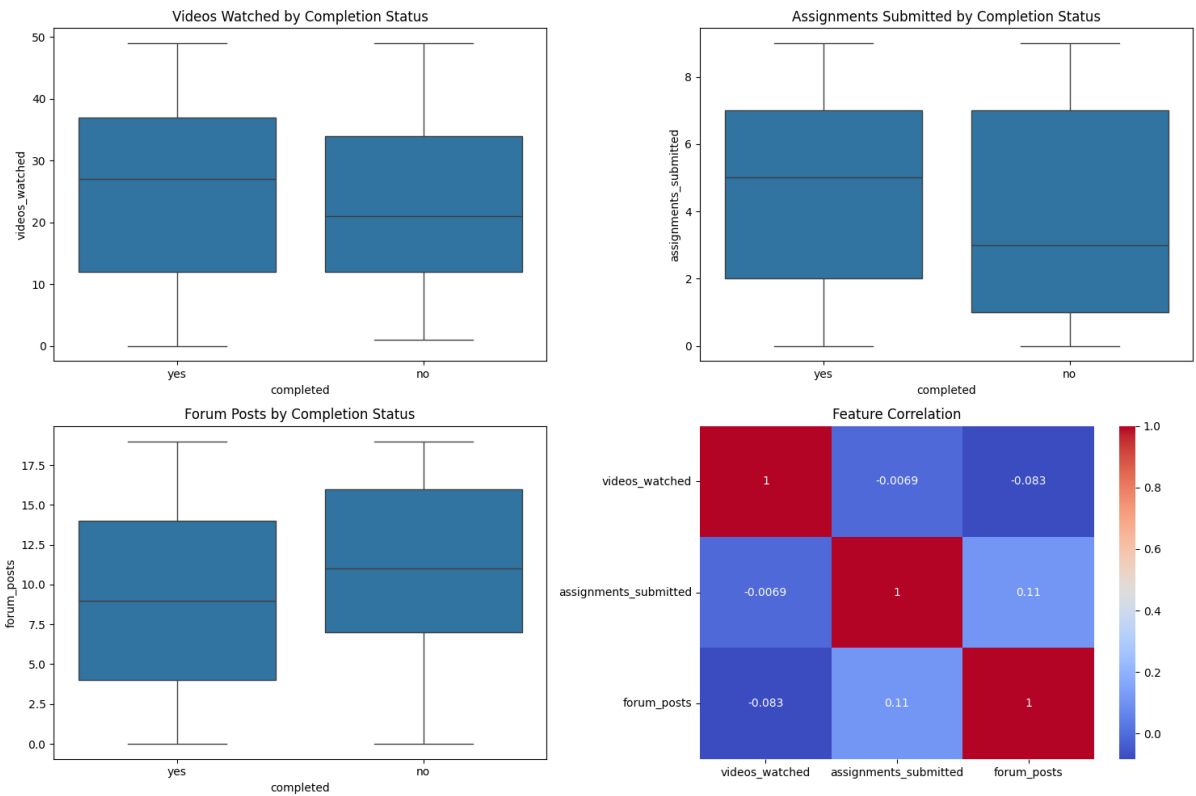
Data summary statistics:

	videos_watched	assignments_submitted	forum_posts
count	100.000000	100.000000	100.000000
mean	23.880000	4.280000	9.500000
std	14.234884	2.930439	5.838742
min	0.000000	0.000000	0.000000
25%	12.000000	2.000000	4.000000
50%	24.000000	4.000000	10.000000
75%	36.250000	7.000000	14.000000
max	49.000000	9.000000	19.000000

2. Class distribution along with visual representation



3. Visual Representation



4. Accuracy checking measures

completed

Logistic Regression Cross-Validation Accuracy: 0.5857 \pm 0.0948
Logistic Regression Test Accuracy: 0.5333

Classification Report:

	precision	recall	f1-score	support
0	0.40	0.55	0.46	11
1	0.67	0.53	0.59	19
accuracy			0.53	30
macro avg	0.53	0.54	0.52	30
weighted avg	0.57	0.53	0.54	30

Random Forest Cross-Validation Accuracy: 0.5143 \pm 0.0833
Random Forest Test Accuracy: 0.5333

Classification Report:

	precision	recall	f1-score	support
0	0.40	0.55	0.46	11
1	0.67	0.53	0.59	19
accuracy			0.53	30
macro avg	0.53	0.54	0.52	30
weighted avg	0.57	0.53	0.54	30

SVM Cross-Validation Accuracy: 0.4571 \pm 0.0969
SVM Test Accuracy: 0.4667


Classification Report:

	precision	recall	f1-score	support
0	0.31	0.36	0.33	11
1	0.59	0.53	0.56	19
accuracy			0.47	30
macro avg	0.45	0.44	0.44	30
weighted avg	0.49	0.47	0.47	30


Best Model: Logistic Regression with accuracy 0.5333

Feature Importance

5. Showing best model after accuracy checking



6. Final Decision boundary output with Model training completion prompt



References & Credits

Dataset

- <https://www.kaggle.com/datasets/rabieelkharoua/predict-online-course-engagement-dataset>
-

External Libraries

- Scikit-learn: Pedregosa, F. et al. (2011). Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research, 12, 2825-2830.
- Pandas: McKinney, W. (2010). Data Structures for Statistical Computing in Python. Proceedings of the 9th Python in Science Conference, 51-56.
- Matplotlib: Hunter, J. D. (2007). Matplotlib: A 2D Graphics Environment. Computing in Science & Engineering, 9(3), 90-95.
- Seaborn: Waskom, M. et al. (2020). Seaborn: Statistical Data Visualization. Journal of Open Source Software, 6(60), 3021.
-

Algorithms & Methodology

- The A* search methodology was adapted from Russell, S., & Norvig, P. (2020). Artificial Intelligence: A Modern Approach (4th ed.). Pearson.
- Random Forest implementation inspired by Breiman, L. (2001). Random Forests. Machine Learning, 45(1), 5-32.
-

Images

- The decision boundary visualization technique is based on examples from the scikit-learn documentation and tutorials.
- Correlation heatmap styling adapted from "Data Visualization with Seaborn" by DataCamp (2023).